

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
import itertools
from functools import partial

# sklearn, model selection
import train_test_split
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, ConfusionMatrixDisplay,

import tensorflow as tf
from tensorflow import keras
```

```
In [21]: %load_ext watermark
%watermark -v -m -p numpy,sklearn,pandas,seaborn,tensorflow
```

Python implementation: CPython  
Python version : 3.8.8  
IPython version : 7.22.0

numpy : 1.23.5  
sklearn : 1.2.0  
pandas : 1.4.3  
seaborn : 0.12.2  
tensorflow: 2.11.0

Compiler : MSC v.1916 64 bit (AMD64)  
OS : Windows  
Release : 10  
Machine : AMD64  
Processor : Intel64 Family 6 Model 140 Stepping 1, GenuineIntel  
CPU cores : 8  
Architecture: 64bit

## Identify a Deep Learning Problem

Neural Networks are the standard for image classification given their complexity. This project is an exploration of different types of NN's learned in the Introduction to Deep Learning course and their ability to classify simple images (fashion-MNIST) compared to simpler machine learning algorithm such as KNN.

The Fashion MNIST dataset is collection of fashion article images, it consists of 70,000 examples in total and is a popular benchmarking dataset of machine learning algorithms as it is more complicated than the regular MNIST dataset.

```
In [185]: fashion_mnist = keras.datasets.fashion_mnist
(X_train, y_train), (X_test, y_test) = fashion_mnist.load_data()
```

```
In [133]: class_names = ["T-shirt/top", "Trouser", "Pullover", "Dress", "Coat",
                        "Sandal", "Shirt", "Sneaker", "Bag", "Ankle boot"]
```

## Exploratory Data Analysis (EDA)

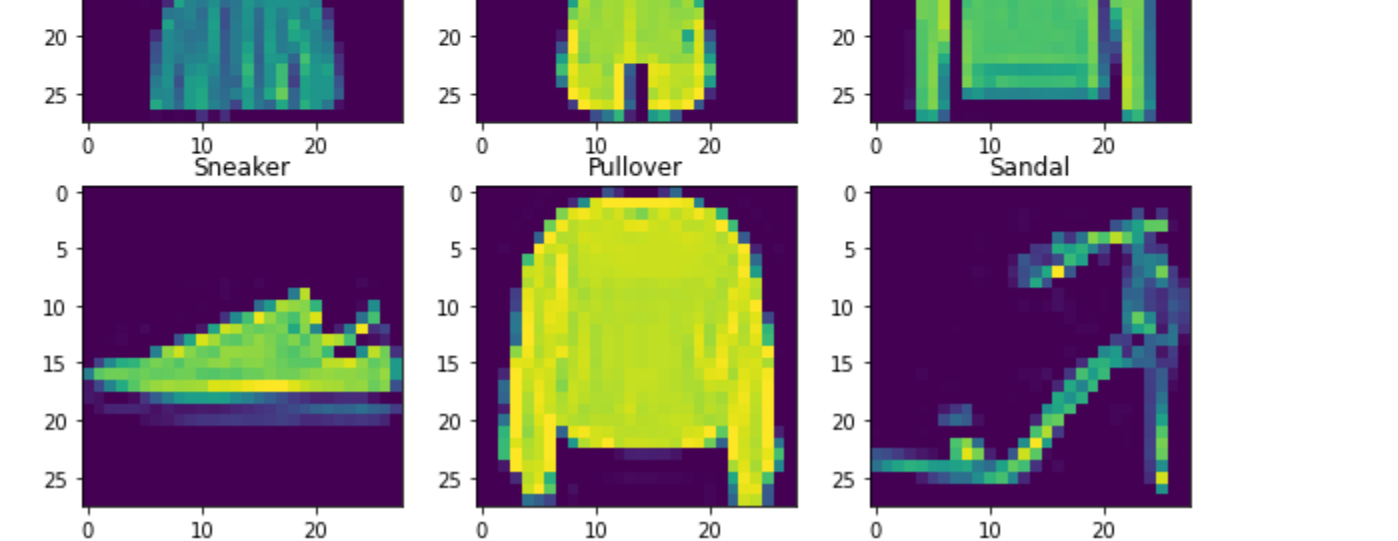
There are a total of 60,000 training and images and a total of 10,000 testing images.The dataframe train\_labels contains the labels for the training images.

The images are 28x28 and are GrayScale images, giving us an input array of 96 by 96 by 1

```
In [134]: print(X_train.shape, X_test.shape)
(60000, 28, 28) (10000, 28, 28)
```

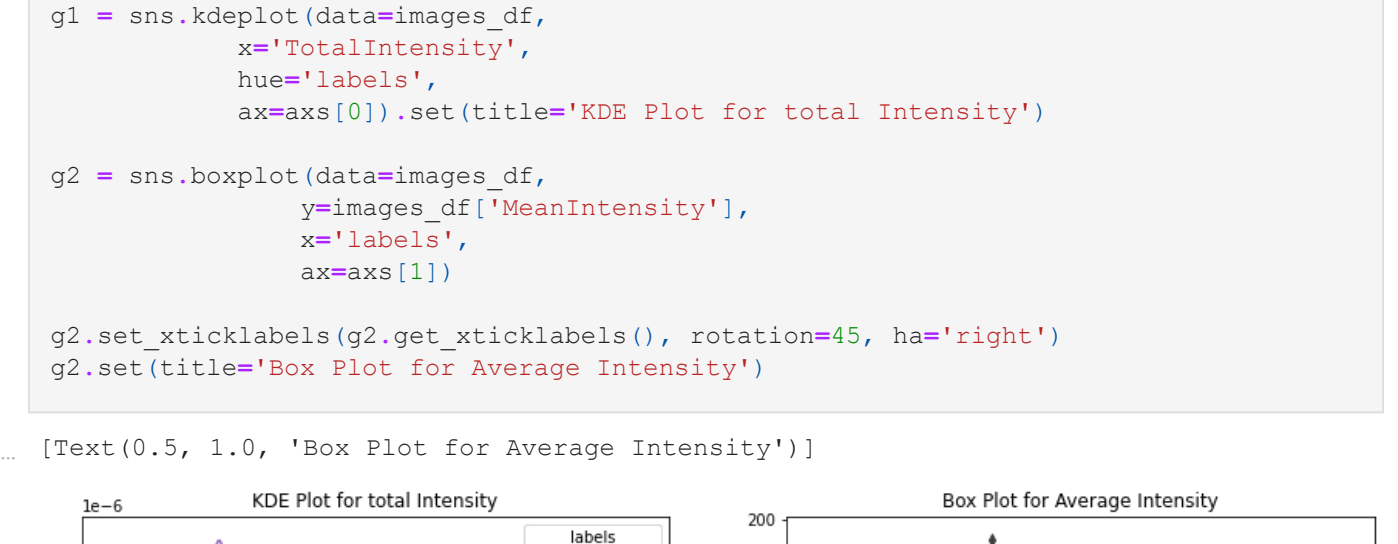
```
In [205]: images_df = pd.DataFrame(np.array(class_names)[y_train], columns=['Labels'])
```

```
In [206]: sns.barplot(data=images_df['Labels'].value_counts().to_frame().reset_index(),
                  x='Labels',
                  y='index')
plt.ylabel('Category')
plt.xlabel('Count of Items')
plt.show()
```



The data is known to be perfectly sampled.

```
In [207]: fig, axes = plt.subplots(3, 3, figsize=(10, 10))
for image, prediction, axis in zip(X_train[:9], np.array(class_names)[y_train[:9]], axes):
    axis.imshow(image)
    axis.set_title(prediction)
```



Single the images are GrayScale, we can look at their sum and average of their intensity values and plot them to get an understanding of how they look like

```
In [228]: images_df['TotalIntensity'] = X_train.reshape((60000, -1)).sum(axis=1)
images_df['MeanIntensity'] = np.mean(X_train.reshape((60000, -1)), axis=1)
```

```
In [239]: figs, axes = plt.subplots(ncols=2,
                              figsize=(15,6))
```

```
g1 = sns.kdeplot(data=images_df,
                  x='TotalIntensity',
                  hue='Labels',
                  ax=axes[0]).set(title='KDE Plot for total Intensity')
g2 = sns.boxplot(data=images_df,
                  x='Labels',
                  y='MeanIntensity',
                  ax=axes[1])
g2.set_title('Box Plot for Average Intensity')
```



## Analysis Using Models

We will analyze three models, a KNN, a MLP and a CNN

### k Nearest Neighbors

We need to reshape the input data for the KNN model (images are 3D, need 2D data)

```
In [240]: knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train.reshape((60000, -1)), y_train)
```

```
Out[240]: KNeighborsClassifier(n_neighbors=3)
```

```
In [241]: knn_predictions = knn.predict(X_test.reshape((10000, -1)))
```

```
In [242]: knn_predictions
```

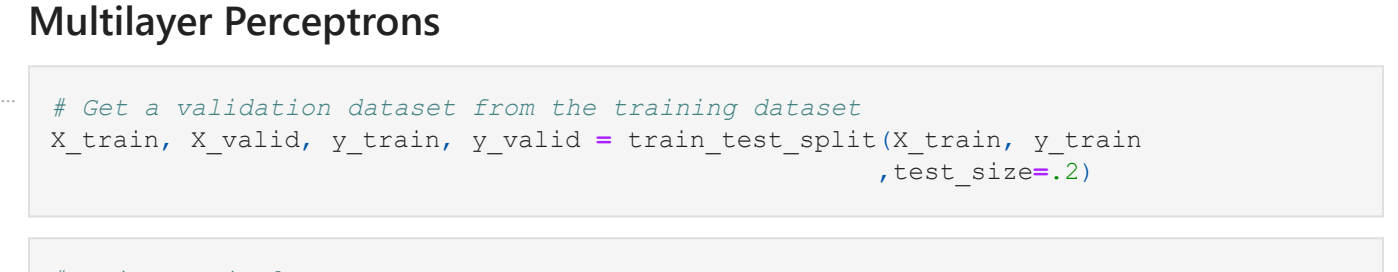
```
Out[242]: array([9, 2, 1, ..., 8, 1, 7], dtype=uint8)
```

### Results

```
In [243]: print("Accuracy ->", accuracy_score(knn_predictions, y_test))
print(classification_report(mlp_predicted_classes, y_test))
```

	precision	recall	f1-score	support
0	0.82	0.81	0.82	1011
1	0.96	0.99	0.98	967
2	0.90	0.62	0.73	1451
3	0.89	0.87	0.88	1017
4	0.64	0.85	0.73	749
5	0.95	0.98	0.97	973
6	0.58	0.71	0.64	826
7	0.96	0.93	0.95	1038
8	0.96	0.98	0.97	973
9	0.95	0.96	0.95	995
accuracy			0.86	10000
macro avg	0.86	0.87	0.86	10000
weighted avg	0.87	0.86	0.86	10000

```
In [198]: cf_matrix = ConfusionMatrixDisplay(confusion_matrix(knn_predictions,
                                                            y_test,
                                                            labels=np.arange(0,10)),
                                          cf_matrix=ax.get_xticklabels(), rotation=45, ha='right')
```



The KNN model had a resulting accuracy of 0.8541, quite impressive.

### Multilayer Perceptrons

```
In [151]: # Get a validation dataset from the training dataset
X_train, X_valid, y_train, y_valid = train_test_split(X_train, y_train,
                                                    test_size=2)
```

```
In [179]: # Using a simple MLP
mlp_model = keras.models.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(400, activation='relu'),
    keras.layers.Dense(200, activation='relu'),
    keras.layers.Dropout(0.2),
    keras.layers.Dense(100, activation='relu'),
    keras.layers.Dropout(0.2),
    keras.layers.Dense(10, activation='softmax')
])
```


```
In [247]: # Get a visual view of the architecture
keras.utils.plot_model(mlp_model, show_shapes=True, show_dtype=True,
                        show_layer_names=True, expand_nested=True,
                        dpi=50, show_layer_activations=True,
                        )
```



```
In [181]: mlp_model.compile(loss="sparse_categorical_crossentropy", # Loss function
                        optimizer="adam", # Gradient Descent
                        metrics=["accuracy"])
mlp_history = mlp_model.fit(X_train, y_train,
                            epochs=50,
                            validation_data=(X_valid, y_valid),
                            batch_size = 10)
```

Epoch 1/50	4600/4800	=====	- 13s 3ms/step - loss: 1.7572 - accuracy: 0.5403 - val_loss: 1.1754 - val_accuracy: 0.5163
Epoch 2/50	4600/4800	=====	- 13s 3ms/step - loss: 1.0797 - accuracy: 0.5426 - val_loss: 1.0644 - val_accuracy: 0.5249
Epoch 3/50	4600/4800	=====	- 12s 2ms/step - loss: 0.8380 - accuracy: 0.6609 - val_loss: 0.7041 - val_accuracy: 0.7230
Epoch 4/50	4600/4800	=====	- 12s 2ms/step - loss: 0.6399 - accuracy: 0.7433 - val_loss: 0.6079 - val_accuracy: 0.7386
Epoch 5/50	4600/4800	=====	- 13s 3ms/step - loss: 0.5993 - accuracy: 0.7555 - val_loss: 0.6046 - val_accuracy: 0.7565
Epoch 6/50	4600/4800	=====	- 13s 3ms/step - loss: 0.5562 - accuracy: 0.7819 - val_loss: 0.6626 - val_accuracy: 0.7825
Epoch 7/50	4600/4800	=====	- 13s 3ms/step - loss: 0.5065 - accuracy: 0.8254 - val_loss: 0.4816 - val_accuracy: 0.8408
Epoch 8/50	4600/4800	=====	- 13s 3ms/step - loss: 0.4561 - accuracy: 0.8461 - val_loss: 0.5200 - val_accuracy: 0.8300
Epoch 9/50	4600/4800	=====	- 12s 3ms/step - loss: 0.4258 - accuracy: 0.8557 - val_loss: 0.4602 - val_accuracy: 0.7230
Epoch 10/50	4600/4800	=====	- 14s 3ms/step - loss: 0.4194 - accuracy: 0.8584 - val_loss: 0.5452 - val_accuracy: 0.8218
Epoch 11/50	4600/4800	=====	- 22s 5ms/step - loss: 0.4121 - accuracy: 0.8625 - val_loss: 0.4386 - val_accuracy: 0.8508
Epoch 12/50	4600/4800	=====	- 15s 3ms/step - loss: 0.3972 - accuracy: 0.7819 - val_loss: 0.4350 - val_accuracy: 0.8722
Epoch 13/50	4600/4800	=====	- 13s 4ms/step - loss: 0.3924 - accuracy: 0.8662 - val_loss: 0.4369 - val_accuracy: 0.8544
Epoch 14/50	4600/4800	=====	- 12s 2ms/step - loss: 0.3899 - accuracy: 0.8686 - val_loss: 0.4679 - val_accuracy: 0.8589
Epoch 15/50	4600/4800	=====	- 13s 3ms/step - loss: 0.3846 - accuracy: 0.8706 - val_loss: 0.4624 - val_accuracy: 0.8454
Epoch 16/50	4600/4800	=====	- 15s 3ms/step - loss: 0.3764 - accuracy: 0.8728 - val_loss: 0.4433 - val_accuracy: 0.8649
Epoch 17/50	4600/4800	=====	- 16s 3ms/step - loss: 0.3776 - accuracy: 0.8730 - val_loss: 0.4516 - val_accuracy: 0.8627
Epoch 18/50	4600/4800	=====	- 14s 3ms/step - loss: 0.3746 - accuracy: 0.8820 - val_loss: 0.4689 - val_accuracy: 0.8450
Epoch 19/50	4600/4800	=====	- 14s 3ms/step - loss: 0.3645 - accuracy: 0.8850 - val_loss: 0.4559 - val_accuracy: 0.8641
Epoch 20/50	4600/4800	=====	- 13s 3ms/step - loss: 0.3649 - accuracy: 0.8767 - val_loss: 0.4234 - val_accuracy: 0.8619
Epoch 21/50	4600/4800	=====	- 16s 4ms/step - loss: 0.3710 - accuracy: 0.8757 - val_loss: 0.5043 - val_accuracy: 0.8593
Epoch 22/50	4600/4800	=====	- 19s 4ms/step - loss: 0.3512 - accuracy: 0.8806 - val_loss: 0.4468 - val_accuracy: 0.8714
Epoch 23/50	4600/4800	=====	- 19s 4ms/step - loss: 0.3594 - accuracy: 0.8789 - val_loss: 0.4416 - val_accuracy: 0.8659
Epoch 24/50	4600/4800	=====	- 18s 3ms/step - loss: 0.3420 - accuracy: 0.8822 - val_loss: 0.4765 - val_accuracy: 0.8669
Epoch 25/50	4600/4800	=====	- 19s 4ms/step - loss: 0.3473 - accuracy: 0.8820 - val_loss: 0.4540 - val_accuracy: 0.8577
Epoch 26/50	4600/4800	=====	- 16s 4ms/step - loss: 0.3499 - accuracy: 0.8828 - val_loss: 0.5024 - val_accuracy: 0.8218
Epoch 27/50	4600/4800	=====	- 19s 4ms/step - loss: 0.3401 - accuracy: 0.8839 - val_loss: 0.5250 - val_accuracy: 0.8598
Epoch 28/50	4600/4800	=====	- 19s 4ms/step - loss: 0.3657 - accuracy: 0.8829 - val_loss: 0.5362 - val_accuracy: 0.8486
Epoch 29/50	4600/4800	=====	- 17s 4ms/step - loss: 0.3415 - accuracy: 0.8864 - val_loss: 0.4996 - val_accuracy: 0.8708
Epoch 30/50	4600/4800	=====	- 16s 3ms/step - loss: 0.3314 - accuracy: 0.8818 - val_loss: 0.6161 - val_accuracy: 0.8534
Epoch 31/50	4600/4800	=====	- 18s 4ms/step - loss: 0.3467 - accuracy: 0.8848 - val_loss: 0.5020 - val_accuracy: 0.8722
Epoch 32/50	4600/4800	=====	- 17s 4ms/step - loss: 0.3375 - accuracy: 0.8867 - val_loss: 0.4780 - val_accuracy: 0.8644
Epoch 33/50	4600/4800	=====	- 16s 3ms/step - loss: 0.3391 - accuracy: 0.8903 - val_loss: 0.8100 - val_accuracy: 0.8691
Epoch 34/50	4600/4800	=====	- 17s 3ms/step - loss: 0.3330 - accuracy: 0.8920 - val_loss: 0.5061 - val_accuracy: 0.8652
Epoch 35/50	4600/4800	=====	- 17s 4ms/step - loss: 0.3192 - accuracy: 0.8934 - val_loss: 0.6005 - val_accuracy: 0.8687
Epoch 36/50	4600/4800	=====	- 17s 4ms/step - loss: 0.3163 - accuracy: 0.8948 - val_loss: 0.5722 - val_accuracy: 0.8652
Epoch 37/50	4600/4800	=====	- 18s 4ms/step - loss: 0.3448 - accuracy: 0.8907 - val_loss: 0.5286 - val_accuracy: 0.8766
Epoch 38/50	4600/4800	=====	- 22s 5ms/step - loss: 0.3328 - accuracy: 0.8926 - val_loss: 0.6711 - val_accuracy: 0.8766
Epoch 39/50	4600/4800	=====	- 17s 4ms/step - loss: 0.3579 - accuracy: 0.8950 - val_loss: 0.6525 - val_accuracy: 0.8642
Epoch 40/50	4600/4800	=====	- 18s 4ms/step - loss: 0.3203 - accuracy: 0.8964 - val_loss: 0.6923 - val_accuracy: 0.8754
Epoch 41/50	4600/4800	=====	- 17s 4ms/step - loss: 0.3252 - accuracy: 0.9051 - val_loss: 0.6087 - val_accuracy: 0.8781
Epoch 42/50	4600/4800	=====	- 17s 3ms/step - loss: 0.3277 - accuracy: 0.8927 - val_loss: 0.6024 - val_accuracy: 0.8678

```
In [182]: pd.DataFrame(mlp_history.history).plot(figsize=(8, 5))
plt.grid(True)
plt.xlabel('Epochs')
plt.gca().set_ylim(0, 1)
plt.title('Accuracy vs Epoch Plot of the MLP Model')
plt.show()
```



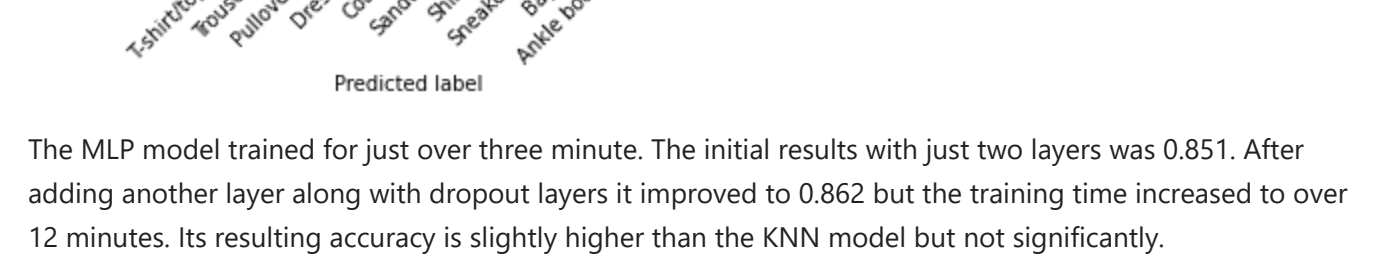
```
In [183]: mlp_predicted_classes = np.argmax(mlp_model.predict(X_test), axis=1)
313/313 [=====] - 1s 2ms/step
```

### Results

```
In [184]: print("Accuracy ->", accuracy_score(mlp_predicted_classes, y_test))
print(classification_report(mlp_predicted_classes, y_test))
```

	precision	recall	f1-score	support
0	0.82	0.81	0.82	1011
1	0.98	1.00	0.99	967
2	0.91	0.76	0.83	1198
3	0.89	0.87	0.88	1017
4	0.64	0.85	0.73	749
5	0.98	0.97	0.98	1019
6	0.58	0.81	0.68	826
7	0.96	0.96	0.96	1005
8	0.98	0.98	0.98	997
9	0.96	0.98	0.97	981
accuracy			0.89	10000
macro avg	0.89	0.90	0.89	10000
weighted avg	0.90	0.89	0.90	10000

```
In [178]: cf_matrix = ConfusionMatrixDisplay(confusion_matrix(mlp_predicted_classes,
                                                            y_test,
                                                            labels=np.arange(0,10)),
                                          cf_matrix=ax.get_xticklabels(), rotation=45, ha='right')
```



The MLP model trained for just over three minute. The initial results with just two layers was 0.851. After adding another layer along with dropout layers it improved to 0.862 but the training time increased to over 12 minutes. Its resulting accuracy is slightly higher than the KNN model but not significantly.

Based on the confusion matrix we can see that the model struggles to identify Shirts, frequently misclassifying them as T-shirts/tops, Pullovers and Coats. The reverse is also a pattern where these are classified as shirts.

### Convolutional Neural Network

```
In [119]: DefaultConv2D = partial(keras.layers.Conv2D, kernel_size=3, activation='relu', padding='same')
```

```
In [120]: cnn_model = keras.models.Sequential([
    DefaultConv2D(filters=64, kernel_size=6, input_shape=(28, 28, 1)),
    keras.layers.MaxPooling2D(pool_size=2), # Pool
    DefaultConv2D(filters=64, kernel_size=6, input_shape=(14, 14, 1)),
    DefaultConv2D(filters=128),
    keras.layers.MaxPooling2D(pool_size=2), # Pool
    DefaultConv2D(filters=256),
    DefaultConv2D(filters=256),
    keras.layers.MaxPooling2D(pool_size=2), # Pool
    keras.layers.Flatten(),
    keras.layers.Dense(units=128, activation='relu'),
    keras.layers.Dropout(0.2),
    keras.layers.Dense(units=64, activation='relu'),
    keras.layers.Dropout(0.2),
    keras.layers.Dense(units=10, activation='softmax'),
])
```

```
In [248]: # Get a visual view of the architecture
keras.utils.plot_model(cnn_model, show_shapes=True, show_dtype=True,
                        show_layer_names=True, expand_nested=True,
                        dpi=20, show_layer_activations=True,
                        )
```

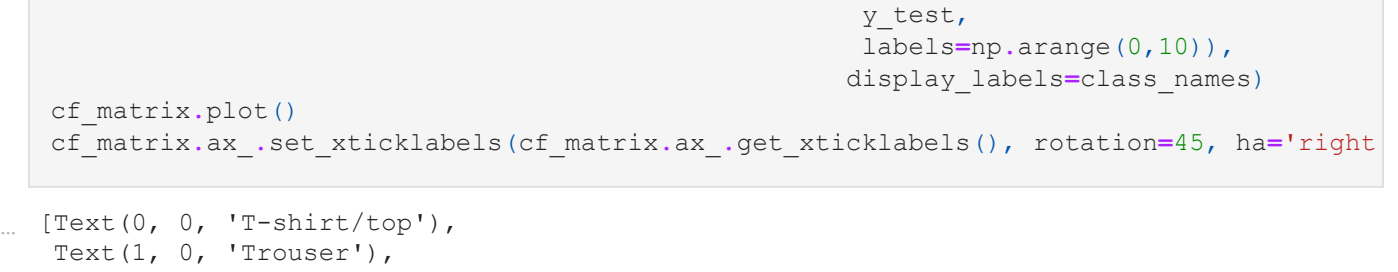


```
Out[248]:
```

```
In [123]: cnn_model.compile(loss="sparse_categorical_crossentropy", # Loss function
                        optimizer="adam", # Gradient Descent
                        metrics=["accuracy"])
cnn_model_history = cnn_model.fit(X_train, y_train,
                                  epochs=20,
                                  validation_data=(X_valid, y_valid))
```

Epoch 1/20	1500/1500	=====	- 210s 139ms/step - loss: 0.5960 - accuracy: 0.7924 - val_loss: 0.3620 - val_accuracy: 0.8450
Epoch 2/20	1500/1500	=====	- 245s 163ms/step - loss: 0.3777 - accuracy: 0.8656 - val_loss: 0.3395 - val_accuracy: 0.8697
Epoch 3/20	1500/1500	=====	- 184s 189ms/step - loss: 0.3362 - accuracy: 0.8868 - val_loss: 0.3184 - val_accuracy: 0.8843
Epoch 4/20	1500/1500	=====	- 275s 188ms/step - loss: 0.3105 - accuracy: 0.8948 - val_loss: 0.2984 - val_accuracy: 0.8921
Epoch 5/20	1500/1500	=====	- 269s 179ms/step - loss: 0.2943 - accuracy: 0.8975 - val_loss: 0.2655 - val_accuracy: 0.9028
Epoch 6/20	1500/1500	=====	- 238s 172ms/step - loss: 0.2754 - accuracy: 0.9010 - val_loss: 0.2897 - val_accuracy: 0.8908
Epoch 7/20	1500/1500	=====	- 240s 180ms/step - loss: 0.2617 - accuracy: 0.9061 - val_loss: 0.2796 - val_accuracy: 0.9020
Epoch 8/20	1500/1500	=====	- 266s 177ms/step - loss: 0.2560 - accuracy: 0.9084 - val_loss: 0.2794 - val_accuracy: 0.9011
Epoch 9/20	1500/1500	=====	- 278s 185ms/step - loss: 0.2401 - accuracy: 0.9147 - val_loss: 0.2847 - val_accuracy: 0.9008
Epoch 10/20	1500/1500	=====	- 286s 190ms/step - loss: 0.2448 - accuracy: 0.9151 - val_loss: 0.2830 - val_accuracy: 0.9039
Epoch 11/20	1500/1500	=====	- 271s 181ms/step - loss: 0.2308 - accuracy: 0.9186 - val_loss: 0.2881 - val_accuracy: 0.9035
Epoch 12/20	1500/1500	=====	- 248s 166ms/step - loss: 0.2405 - accuracy: 0.9163 - val_loss: 0.2934 - val_accuracy: 0.9048
Epoch 13/20	1500/1500	=====	- 248s 166ms/step - loss: 0.2337 - accuracy: 0.9184 - val_loss: 0.2707 - val_accuracy: 0.9039
Epoch 14/20	1500/1500	=====	- 241s 160ms/step - loss: 0.2293 - accuracy: 0.9209 - val_loss: 0.2821 - val_accuracy: 0.9012
Epoch 15/20	1500/1500	=====	- 248s 166ms/step - loss: 0.2253 - accuracy: 0.9214 - val_loss: 0.2794 - val_accuracy: 0.9101
Epoch 16/20	1500/1500	=====	- 238s 159ms/step - loss: 0.2379 - accuracy: 0.9181 - val_loss: 0.3095 - val_accuracy: 0.9040
Epoch 17/20	1500/1500	=====	- 238s 154ms/step - loss: 0.2186 - accuracy: 0.9232 - val_loss: 0.3042 - val_accuracy: 0.9118
Epoch 18/20	1500/1500	=====	- 246s 164ms/step - loss: 0.2311 - accuracy: 0.9206 - val_loss: 0.3155 - val_accuracy: 0.8980

```
In [125]: pd.DataFrame(cnn_model_history.history).plot(figsize=(8, 5))
plt.grid(True)
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.gca().set_ylim(0, 1)
plt.title('Accuracy vs Epoch Plot of the CNN Model')
plt.show()
```



```
In [126]: cnn_model_classes = np.argmax(cnn_model.predict(X_test), axis=1)
313/313 [=====] - 24s 75ms/step
```

### Results

```
In [127]: print("Accuracy ->", accuracy_score(cnn_model_classes, y_test))
print(classification_report(cnn_model_classes, y_test))
```

True label		Predicted label										Support		
		Coat	Coat	Sandals	Shirts	Sneakers	Bag	Ankle boot	Coat	Sandals	Shirts		Sneakers	Bag
Coat	1	2	31	41	788	0	70	0	1	0	0	0	0	0
Sandals	1	0	0	1	0	991	0	19	1	12	0	0	0	0
Shirts	80	0	25	16	48	0	559	0	5	0	0	0	0	0
Sneakers	0	0	0	0	0	11	1	962	1	30	0	0	0	0
Bag	7	0	0	1	1	0	6	1	981	0	0	0	0	0
Ankle boot	0	0	0	1	0	4	0	18	1	957	0	0	0	0